

Machine Learning Engineer Nanodegree Capstone Proposal

Sebastian Schmitz

March 7, 2018

1 Domain background

Ever since Tesauro's backgammon agent TD-Gammon [6] proved that computer programs can perform as good as humans in certain games, machine learning algorithms have been applied in numerous games. This led to advances on both sides: The machine learning community could refine and improve their algorithms and the gaming communities learned new ways to play.

TD-Gammon used openings that had wrongfully been regarded as inferior by the leading players in the world. After a few tournaments where players tested the more conservative moves chosen by TD-Gammon, they became the near-universal choice, because they are - in fact - superior[4].

One of the pinnacles of this development was reached when the AlphaGo computer program defeated the European Go champion Fan Hui on a full 19x19 board[2]. Until then, Go[1] had been regarded as too complex to be mastered by an artificial intelligence.

Fueled by this development artificial intelligence programs trained through various forms of reinforcement learning[3] have learned to fly helicopters [5], play soccer [?] and optimize stock trade execution [?].

Inspired by the ability of letting an agent learn how to thrive in an environment without being told what to do, I developed a reinforcement learning agent which solved a version of the game Rush Hour¹. The game consists of a 2d board with wooden pieces on it. The goal is to bring one of the pieces into a special position only by moving the blocks left, right, up or down. My approach was able to find a solution that was a lot shorter than the solution that came with the game.

I now want to build upon this and move from a discrete problem to a continuous one.

2 Problem statement

As an exercise to learn C++ and take a first peak into game programming, I programmed a 2d space game at the beginning of my studies in 2010, see figure ?? . The goal of the game is to reach a score as high as possible by surviving and shooting down asteroids. These asteroids spawn at random positions on the top of the screen and can either hit the player's ship making him lose a life or hit the off-screen space station which the player is supposed to protect. The space station gets repaired over time, but the player will lose if its health drops below 0. It might therefore be wise to crash the ship into an asteroid if it cannot be stopped otherwise and the space station is at low health. In other situations it could be beneficial to let the asteroids hit the space station, since its health bar is still full. Additionally, the amount of shots the player can fire at a time is limited, so it's not feasible to just fire at will.

The task of this project is therefore to let a machine learning algorithm figure out how to achieve the highest score. In order not to bias the learning agent, no human-devised strategies will be programmed into the approach. This will lead to an algorithm that potentially uses methods a human would not have thought of and that can cope with any of the random situations, the game throws at it.

¹[https://en.wikipedia.org/wiki/Rush_Hour_\(board_game\)](https://en.wikipedia.org/wiki/Rush_Hour_(board_game))




Figure 1: 2d space game

3 Dataset and Inputs

Since this is a simulation on a custom game there is no dataset to use. The inputs the algorithm can make use of are the following:

- **Lives left:** Every time the space ship gets hit by an asteroid, the player will lose a life and eventually the entire game if no lives are left.
- **Space station health:** Every time an asteroid is allowed to cross the screen, the space station will lose a set amount of health and regenerate it at a fixed rate until it's back at 100 again. If the space station health drops below 0, the player will lose the game.
- **Asteroids:** Set of x- and y-coordinates describing the positions of the asteroids
- **Ship position:** The ship can move left and right on the lower corner of the screen
- **Shots left:** Only a certain amount of shots can be fired at a time.
- **Shot positions:** Set of x- and y-coordinates of the shots the player has already fired.

4 Solution statement

In the MLND lectures we learned about three fields of machine learning: Supervised learning, unsupervised learning and reinforcement learning. The proposed problem does not include large datasets that we can mine for information which we would need to apply (un)supervised learning methods. It is rather an environment in which an algorithm has to learn over time, adapt to so far unseen situations and generalize the experiences made in order to be able to handle as many problems that can arise within this environment as possible.

The solution will therefore leverage the strength of algorithms coming from the field of reinforcement learning: The system will start with no knowledge about the problem whatsoever, but it will perceive the state of the environment and will be able to derive a set of actions from this state. At first, the actions taken will be randomly chosen: The environment transitions into another state and the agent will immediately receive feedback on whether the choice of action was good or not. From these feedbacks the algorithm will over time create a graph modelling the environment: The nodes will be the states, the edges and their weights will be the actions. Since the goal is to learn how to behave, the level of randomness when choosing actions will be decreased by a small amount with every run. This leads to a transition in the behavior: In the beginning, the agent will explore the environment and after some time start exploiting its knowledge to ensure that the best solution was found.

The result will be an agent that has explored a big portion of the state space and can play the game reasonably well, ideally better than a human player.

5 Benchmark Model

Since the agent will train on a custom game, the choice of benchmark models and solutions to compare the performance to is sparse:

6 Evaluation Metrics

7 Project Design

References

- [1] Go (game). [https://en.wikipedia.org/wiki/Go_\(game\)](https://en.wikipedia.org/wiki/Go_(game)). Accessed: 2017-11-08.
- [2] In a huge breakthrough, google's ai beats a top player at the game of go. <https://www.wired.com/2016/01/in-a-huge-breakthrough-googles-ai-beats-a-top-player-at-the-game-of-go/>. Accessed: 2017-11-08.
- [3] Reinforcement learning. https://en.wikipedia.org/wiki/Reinforcement_learning. Accessed: 2017-11-11.
- [4] Temporal difference learning and td-gammon. <http://www.bkgm.com/articles/tesauro/td1.html>. Accessed: 2017-11-08.

- [5] Pieter Abbeel, Adam Coates, Morgan Quigley, and Andrew Y. Ng. An application of reinforcement learning to aerobatic helicopter flight. In P. B. Schölkopf, J. C. Platt, and T. Hoffman, editors, *Advances in Neural Information Processing Systems 19*, pages 1–8. MIT Press, 2007.
- [6] Gerald Tesauro. Temporal difference learning and td-gammon. *Commun. ACM*, 38(3):58–68, March 1995.